

Model Development Phase Template

Date	10 JULY 2024
Team ID	SWTID1720000556
Project Title	Predicting Co2 Emission By Countries Using Machine Learning
Maximum Marks	4 Marks

Initial Model Training Code, Model Validation and Evaluation Report

The initial model training code will be showcased in the future through a screenshot. The model validation and evaluation report will include classification reports, accuracy, and confusion matrices for multiple models, presented through respective screenshots.

Initial Model Training Code:

1)

```
linear Regression

[124]: lin_reg = LinearRegression()

[125]: lin_reg.fit(x_train, y_train)

[125]: LinearRegression
LinearRegression()

[180]: one_pred_lin_reg=lin_reg.predict(x_test)
two_pred_lin_reg=lin_reg.predict(x_train)

[171]: mse_lin_reg = mean_squared_error(y_test, y_pred_lin_reg)
print(f'Linear Regression MSE: {mse_lin_reg}')
lin_reg.score(x_test,y_test)

Linear Regression MSE: 2.3664620738212395e+27
[171]: 0.0002686736146634283

[182]: print(lin_reg.score(x_train,y_train))
print(f"mse for train data set", mean_squared_error(y_train,two_pred_lin_reg))

0.00024196666451214988
mse for train data set 2.3388476124039764e+27
```

2)

Ridge Regression

```
[134]: ridge_reg = Ridge()

[135]: x=np.logspace(-4, 4, 50)
      param_distributions = {'alpha': x}

[136]: ridge_random = RandomizedSearchCV(ridge_reg, param_distributions, n_iter=10, random_state=42)
      ridge_random.fit(x_train, y_train)

[136]: > RandomizedSearchCV ⓘ ⓘ
      > best_estimator_: Ridge
      > Ridge ⓘ

[137]: ridgebest_alpha = ridge_random.best_params_
      best_score = ridge_random.best_score_
      print(best_score)
      print(ridgebest_alpha)
      0.0002386836157977523
      {'alpha': 6866.488450042998}

[138]: ridge_random.score(x_test, y_test)

[138]: 0.0002686734707454397

[161]: y_pred3=ridge_random.predict(x_test)

[163]: mse5=mean_squared_error(y_test, y_pred3)
      mse5

[163]: 2.3664620741619075e+27
```

3)

Decision Tree Regressor

```
[139]: tree_reg = DecisionTreeRegressor()
param_distributions = {
    'max_depth': randint(1, 20),
    'min_samples_split': randint(2, 20),
    'min_samples_leaf': randint(1, 20),
    'max_features': ['auto', 'sqrt', 'log2', None]
}
```

```
[140]: tree_random = RandomizedSearchCV(tree_reg, param_distributions, n_iter=100, cv=5, random_state=42, n_jobs=-1)
tree_random.fit(x_train, y_train)
```

```
2.02224910e-01 1.57061184e-01 3.03250071e-04 7.84325944e-02
4.06398784e-01 6.08873859e-02 nan nan
1.84949540e-03 2.50626508e-01 1.84949540e-03 1.53368819e-01
3.28675273e-04 1.22658713e-01 3.34560757e-01 2.31624067e-03
6.72816772e-01 1.27770142e-01 1.51473417e-01 3.57935091e-01
3.90789107e-01 3.72512052e-01 1.33312224e-01 1.02328613e-03
1.16241165e-01 nan nan nan
1.07312609e-01 5.80033914e-01 7.23929234e-01 nan
1.13488537e-01 nan 1.05792997e-01 1.43884263e-01]
warnings.warn()
```

```
[140]: RandomizedSearchCV
  > best_estimator_: DecisionTreeRegressor
    > DecisionTreeRegressor
```

```
[189]: test_pred_tree = tree_random.predict(x_test)
train_pred_tree = tree_random.predict(x_train)
```

```
[141]: best_params = tree_random.best_params_
best_score = tree_random.best_score_
print(best_params)
print(best_score) # for train data Score
```

```
{'max_depth': 19, 'max_features': None, 'min_samples_leaf': 5, 'min_samples_split': 10}
0.72392923376112
```

```
[2]: tree_random.score(x_test, y_test)
```

```
[2]: 0.7593560151530763
```

```
[1]: mse1 = mean_squared_error(test_pred_tree, y_test)
mse2 = mean_squared_error(train_pred_tree, y_train)
```

```
[3]: print(mse2) # train
print(mse1) # test
```

```
3.3566794752688856e+26
5.696279074223585e+26
```

4)

XGBoost regression

```
xgb_reg = XGBRegressor(objective='reg:squarederror', n_estimators=100, learning_rate=0.1)
xgb_reg.fit(x_train, y_train)
```

```
XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.1, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=100, n_jobs=None,
              num_parallel_tree=None, random_state=None, ...)
```

```
y_pred_xgb = xgb_reg.predict(x_test)
```

```
mse_xgb = mean_squared_error(y_test, y_pred_xgb)
print(f'XGBoost Regression MSE: {mse_xgb}')
print(xgb_reg.score(x_train,y_train))
print(xgb_reg.score(x_test,y_test))

XGBoost Regression MSE: 1.7127083409170892e+27
0.3248889392152211
0.2764523039190644
```

5)

Random Forest Regression

```
model=RandomForestRegressor(n_estimators=10,random_state=52,n_jobs=-1)
```

```
model.fit(x_train,y_train)
```

```
RandomForestRegressor
RandomForestRegressor(n_estimators=10, n_jobs=-1, random_state=52)
```

```
y_pred_random =model.predict(x_test)
train_pred=model.predict(x_train)
mse_random = mean_squared_error(y_test, y_pred_random)
print(f'Random Forest Regression MSE: {mse_random}')
model.score(x_train,y_train)
```

```
Random Forest Regression MSE: 1.700880772252512e+26
0.9830788955337789
```

```
mse_random_train=mean_squared_error(y_train,train_pred)
model.score(x_test,y_test)
```

```
0.9281448957378994
```

```
# as we can see Random Forest Regressor has better score and mse
# so, we are continuing with Random Forest Regressor as the model
```

Model Validation and Evaluation Report:

Model	Classification Report (Adjusted R2 Score)	R2 score	Mean Squared Error (MSE) And Mean Absolute Error
Linear Regres sion	<pre>adjusted_r_squared = 1 - (1 - r_squared) * (n - 1) / (n - k - 1) print(adjusted_r_squared) 0.00023312926973717563</pre>	0.00024	<pre>: mae_linear=mean_absolute_error(y_test,one_pred_lin_reg) print(f"mean absolute error for linear_regression is :{mae_linear}") mean absolute error for linear_regression is :2179390902451.5435</pre> <hr/> <pre>: mse_lin_reg = mean_squared_error(y_test, one_pred_lin_reg) print(f'Linear Regression MSE: {mse_lin_reg}') lin_reg.score(x_test,y_test) Linear Regression MSE: 2.3664620738212395e+27</pre>
Ridge Regres sion	<pre>adjusted_r_squared1 = 1 - (1 - r_squared1) * (n - 1) / (n - k - 1) print(f"adjusted R squared value is:{adjusted_r_squared1}") adjusted R squared value is:0.00023312926972629544</pre>	0.00026	<pre>: mae_ridge=mean_absolute_error(y_test,y_pred3) print(f"mean absolute error for ridge_regression is :{mae_ridge}") mean absolute error for ridge_regression is :2179389203351.5894</pre> <hr/> <pre>mse5=mean_squared_error(y_test,y_pred3) mse5 2.3664620741619075e+27</pre>
Decisi on Tree Regres sor	<pre>r_squared2=tree_random.score(x_train,y_train) print(f"R2 score value is: {r_squared2}") adjusted_r_squared2 = 1 - (1 - r_squared2) * (n - 1) / (n - k - 1) print('adjusted R2 score value is :',adjusted_r_squared2) R2 score value is: 0.8565162068304613 adjusted R2 score value is : 0.8565149385006436</pre>	0.75935	<pre>mse1=mean_squared_error(test_pred_tree,y_test) mse2=mean_squared_error(train_pred_tree,y_train) print(mse2) # train print(mse1) # test 3.356679475268885e+26 5.696279458767446e+26</pre> <hr/> <pre>mae_DTree=mean_absolute_error(y_test,test_pred_tree) print(f"mean absolute error for linear_regression is :{mae_DTree}") mean absolute error for linear_regression is :517578714625.5952</pre>

<p>XG boost Regres sion</p>	<pre> r_squared3=xgb_reg.score(x_train,y_train) adjusted_r_squared3 = 1 - (1 - r_squared3) * (n - 1) / (n - k - 1) print(f"the adjusted R2 value is:{adjusted_r_squared3}") the adjusted R2 value is:0.3248829715482836 </pre>	<p>0.32488</p>	<pre> mae_XGB=mean_absolute_error(y_test,y_pred_xgb) print(f"mean absolute error for linear_regression is :{mae_XGB}") mean absolute error for linear_regression is :2051870348116.5146 mse_xgb = mean_squared_error(y_test, y_pred_xgb) print(f'XGBoost Regression MSE: {mse_xgb}') print(xgb_reg.score(x_train,y_train)) print(xgb_reg.score(x_test,y_test)) XGBoost Regression MSE: 1.7127083409170892e+27 0.3248889392152211 0.2764523039190644 </pre>
<p>Rando m Forest Regres sor</p>	<pre> adjusted_r_squared4 = 1 - (1 - score1) * (n - 1) / (n - k - 1) print(f"adjusted_r_squared value is :{adjusted_r_squared4}") adjusted_r_squared value is :0.9830787459591066 </pre>	<p>0.98307</p>	<pre> y_pred_random =model.predict(x_test) train_pred=model.predict(x_train) mse_random = mean_squared_error(y_test, y_pred_random) print(f'Random Forest Regression MSE: {mse_random}') score1=model.score(x_train,y_train) Random Forest Regression MSE: 1.7008807722525123e+26 mae_Random=mean_absolute_error(y_test,y_pred_random) print(f"mean absolute error for linear_regression is :{mae_Random}") mean absolute error for linear_regression is :188752553779.6081 </pre>